

Weekly Status Report (Jan 26 - Feb 2)

Yuetong Liu

02/02/2020

Abstract

Dynamic Programming is a collection of algorithms that can be used to compute optimal policies given a perfect model of the environment as a Markov decision process. It actually consists of two different versions of how it can be implemented which are policy iteration and value iteration. This report concludes the implementation of both algorithms and performance test results in the chess game environment.

1 Accomplishments (Summary)

Policy iteration and value iteration are applied to four agents: king, knight, bishop, and rook. For all agents, both two methods significantly shorten the paths between 2 squares, while policy iteration yields more stable results but value iteration takes less time to converge.

2 Introduction

Tackling chess is challenging because of its huge state space. Therefore I start with finding the shortest path between 2 squares on a chessboard. This problem has a small state space, therefore allows us to tackle this with simple RL algorithms. Moreover, moving chess effectively is an intermediate step to solve the chess problem. This time I used two Dynamic Programming technics - policy iteration and value iteration to solve the problem. As model-based reinforcement learning methods, policy iteration/value iteration tries to understand the chess game and create a model to represent it.

3 Mathematical Development

Below is the algorithm for policy iteration, where ϵ is a small positive number determining the accuracy of the estimation. Section 4 contains python code for this algorithm and its transform(value iteration).

- Initialize policy π_0 randomly
- Iterate until π_i converge to π^*

1. Policy Evaluation:

$$\delta = 0$$

while $\delta \geq \epsilon$:

$$v = V(s)$$

$$V(s) = \sum P(s, \pi(s), s')[r(s, \pi(s), s') + \gamma V(s')] \quad (*)$$

$$\delta = \max(\delta, v - V(s))$$

2. Policy Improvement:

$$\pi_{i+1}(s) = \operatorname{argmax} \sum P(s, a, s') [r(s, a, s') + \gamma V^*(s')] (**)$$

if $\pi_{i+1}(s) \neq \pi_i(s)$: go back to Policy Evaluation

4 Experimental Approach and Results

4.1 Experimental Approach

This part includes python code for algorithms mention in section 3 and the comments on codes are highlighted in green.

4.1.1 State Evaluation

A state (s) is as valuable (V) as the successor state (s') plus the reward (R) for going from s to s'. Since there can be mulitple actions (a) and multiple successor states they are summed and weighted by their probability (pi). In a non-deterministic environment, a given action could result in multiple successor states. We don't have to take this into account for this problem because move chess is a deterministic game. Successor state values are discounted with discount factor (gamma) that varies between 0 and 1.

```
1 def evaluate_state(self, state, gamma=0.9, synchronous=True):
2     Calculates the value of a state based on the successor states and
        the immediate rewards.Args:state: tuple of 2 integers 0-7
        representing the stategamma: float, discount factorsynchronous:
        BooleanReturns: The expected value of the state under the
        current policy.
3     greedy_action_value = np.max(self.agent.policy[state[0],
        state[1], :]) # get max action value at give state
4     greedy_indices = [i for i, a in enumerate(self.agent.policy
        [state[0], state[1], :]) if
5         a == greedy_action_value] # list the
        index of all actions with max action
```

```

                                value
6      prob = 1 / len(greedy_indices) # probability of an action
                                occurring
7      state_value = 0 # set  $V(S) = 0$ 
8      for i in greedy_indices: # for all actions with max action
                                value
9          self.env.state = state # reset state to the one being
                                evaluated
10         reward, episode_end = self.env.step(self.agent.
                                action_space[i]) # get reward
11         if synchronous: # get successor state value
12             successor_state_value = self.agent.
                                value_function_prev[self.env.state]
13             # if synchronous, successor state value is in the
                                same iteration of policy evaluation
14         else:
15             # if not, successor state value could be previous
                                or the current value function, or combined
16             successor_state_value = self.agent.value_function[
                                self.env.state] # otherwise
17         state_value += (prob * (
18             reward + gamma * successor_state_value)) # sum
                                up rewards and discounted successor state
                                value in equation(*)
19     return state_value

```

4.1.2 Policy Evaluation

python Apply state evaluation to all states. *def* *evaluate_policy*(*self*, *gamma* = 0.9, *synchronous* = *True*) :
self.*agent*.*value_function_prev* = *self*.*agent*.*value_function*.*copy*() *For* *synchronous* *updates* *for* *row* *in* *range*(*self*.*agent*.*value_function*.*shape*[0]) :
in *each* *row* *for* *col* *in* *range*(*self*.*agent*.*value_function*.*shape*[1]) : *in* *each* *column* *self*.*agent*.*value_function*[*row*, *col*] =
self.*evaluate_state*((*row*, *col*), *gamma* = *gamma*, *in* *each* *action*, *update* *value* *functions* *synchronous* =
synchronous)

4.1.3 Policy Improvement

Policy Improvement is the act of making the policy greedy with respect to the value function.

```
1 def improve_policy(self):
2     Finds the greedy policy w.r.t. the current value function
3
4     self.agent.policy_prev = self.agent.policy.copy() # get pi(
5         i)
6     for row in range(self.agent.action_function.shape[0]): #
7         for each row
8         for col in range(self.agent.action_function.shape[1]):
9             # for each column
10            for action in range(self.agent.action_function.
                shape[2]): # for each action
                self.env.state = (row, col) # reset state to
                    the one being evaluated
                reward, episode_end = self.env.step(self.agent.
                    action_space[action]) # make step based on
                        action
                successor_state_value = 0 if episode_end else
                    self.agent.value_function[self.env.state] #
```

```

11         update successor state value
12         self.agent.policy[row, col, action] = reward +
13             successor_state_value # get pi(i+1) using
14             (**)
15
16     max_policy_value = np.max(self.agent.policy[row,
17         col, :]) # get max policy value at given state
18     max_indices = [i for i, a in enumerate(self.agent.
19         policy[row, col, :]) if a == max_policy_value] #
20         get the index of policy with max policy value
21     for idx in max_indices:
22         self.agent.policy[row, col, idx] = 1 # 1 if the
23             policy value is max, otherwise 0

```

4.1.4 Policy Iteration

Policy Iteration finds the optimal policy by doing policy evaluation and policy improvement until the policy is stable.

```

1 def policy_iteration(self, eps=0.1, gamma=0.9, iteration=1, k=32,
2     synchronous=True):
3     Finds the optimal policy
4     Args:eps: float, exploration rate
5     gamma: float, discount factor
6     iteration: the iteration number
7     k: (int) maximum amount of policy evaluation
8     iterationssynchronous: (Boolean) whether to use synchronous or asynchronous
9     back-upsReturns:
10     policy_stable = True
11     print(iteration:, iteration, )
12     print( policy:)
13     self.visualize_policy()

```

```

7
8     # Evaluate Policy
9     print()
10    value_delta_max = 0
11    for _ in range(k): # max number of iteration for policy
        evaluation
12        self.evaluate_policy(gamma=gamma, synchronous=
            synchronous) # update V(s)
13    value_delta = np.max(np.abs(self.agent.
        value_function_prev - self.agent.value_function)) #
        get delta
14    value_delta_max = value_delta
15    if value_delta_max < eps: # terminate the loop if delta
        is small
16        break
17    print(Value function for this policy:)
18    print(self.agent.value_function.round().astype(int))
19    action_function_prev = self.agent.action_function.copy()
20
21    # Improve Policy
22    print( Improving policy:)
23    self.improve_policy()
24
25    # Check if Stable
26    policy_stable = self.agent.compare_policies() < 1 #check if
        policy function is similar with previous iteration
27    print(policy diff:, policy_stable)

```

```

28
29     if not policy_stable and iteration < 1000: # if policy is
        not stable, restart the whole function
30         iteration += 1
31         self.policy_iteration(iteration=iteration)
32     elif policy_stable:
33         print(Optimal policy found in, iteration, steps of policy
            evaluation)
34     else:
35         print(failed to converge.)

```

4.1.5 Value Iteration

Theory Value iteration is nothing more than a simple parameter modification to policy iteration. Remember that policy iteration consists of policy evaluation and policy improvement. The policy evaluation step does not necessarily have to be repeated until convergence before we improve our policy. If we use only 1 iteration instead we call it value iteration.

4.2 Results

I selected four agents for training using both policy iteration and value iteration.

- King can move exactly one square horizontally, vertically, or diagonally.
- Rook can move any number of vacant squares vertically or horizontally.
- Bishop can move any number of vacant squares in any diagonal direction.
- Knight can move one square along any rank or file and then at an angle.

Each scenario is repeated 50 times to reduce bias. Here is an example using value iteration to train the agent “King” 50 times. The output is the number of steps the King take from (0,0) to (5, 7)

```

1  eva3 = [] # a empty list to store number of step to get to desired
    position
2  for i in range(50): # repeat 50 times
3      p = Piece(piece='king') # select a chess agent (knight, bishop or
        rook)
4      env = Board() # 8*8 chess board
5      r = Reinforce(p,env) # place king on the board at (0,0) by
        default
6      r.policy_iteration(k=1) # when k=1, it's value iteration
7      states, actions, rewards = r.play_episode(state = (0,0),
8          max_steps=1e3,
9          epsilon=0.1) # move chess until it gets to (5,7)
        by default
10
11     eva3.append(len(actions)) # add number of steps to the list
12 eva3 # print list

```

Output: [11, 8, 9, 10, 8, 8, 8, 8, 9, 8, 8, 8, 8, 10, 12, 10, 9, 8, 8, 8, 8, 10, 11, 8, 8, 8, 11, 8, 12, 8, 12, 8, 8, 8, 8, 10, 8, 8, 8, 8, 8, 10, 10, 10, 8, 8, 10, 9, 8, 12, 9]

The tables show the mean and variance of the total number of steps that an agent takes from (0,0) to (7,5).

Method	King	Knight	Bishop	Rook
None	213.2	169.38	157.36	141.82
Policy Iteration	8.96	5.58	3.28	3.22
Value Iteration	9.10	5.42	3.28	3.22

Table 1: Mean of NO. steps

None	24388.48	22948.03	30279.23	17110.63
Policy Iteration	1.52	1.47	0.36	0.29
Value Iteration	2.65	1.47	0.32	0.21

Table 2: Variance of NO. steps

According to the result table, the mean and variance for each agent reduce significantly after policy iteration/value iteration, while the difference between policy iteration and value iteration is not significant.

5 Next Steps

It is expensive to train them since they required iterations to converge. Moreover, policy evaluation calculates the state value by backing up the successor state values and the transition probabilities to those states. However, these probabilities are usually unknown. To work with unknown environments some model-free techniques such as Monte Carlo Control could be an alternative approach.

References

- [1] Richard S. Sutton and Andrew G. Barto. 2018. Reinforcement Learning: An Introduction. A Bradford Book, Cambridge, MA, USA.
- [2] Groen, A. 2019. Reinforcement Learning Chess. GitHub repository, <https://github.com/arjangroen/RLC>