

# RL\_Deep\_Q\_Week\_4

November 10, 2020

## 1 Import Package

```
[1]: import pyvirtualdisplay
import matplotlib.pyplot as plt
import tensorflow as tf
from tf_agents.agents.dqn import dqn_agent
from tf_agents.environments import suite_gym
from tf_agents.environments import tf_py_environment
from tf_agents.networks import q_network
from tf_agents.utils import common
from tf_agents.replay_buffers import tf_uniform_replay_buffer
from tf_agents.trajectories import trajectory
import PIL.Image
from tf_agents.policies import random_tf_policy

[2]: tf.compat.v1.enable_v2_behavior()
# Set up a virtual display for rendering OpenAI gym environments.
#display = pyvirtualdisplay.Display(visible=0, size=(1400, 900)).start()
```

## 2 Hyperparameters

```
[3]: num_iterations = 20000 # @param {type:"integer"}

initial_collect_steps = 100 # @param {type:"integer"}
collect_steps_per_iteration = 1 # @param {type:"integer"}
replay_buffer_max_length = 100000 # @param {type:"integer"}

batch_size = 64 # @param {type:"integer"}
learning_rate = 1e-3 # @param {type:"number"}
log_interval = 200 # @param {type:"integer"}

num_eval_episodes = 10 # @param {type:"integer"}
eval_interval = 1000 # @param {type:"integer"}
```

### 3 Environment

```
[4]: env_name = 'CartPole-v0'
     env = suite_gym.load(env_name)

[5]: # Train and Evaluation(Testing)
     train_py_env = suite_gym.load(env_name)
     eval_py_env = suite_gym.load(env_name)

     # Python to TF
     train_env = tf_py_environment.TFPyEnvironment(train_py_env)
     eval_env = tf_py_environment.TFPyEnvironment(eval_py_env)
```

### 4 Agent / Algorithm

$$Y(s,a,r,s) = r + \max_a Q(s,a)$$

$$L() = (s,a,r,s) \sim U(D) [(Y(s,a,r,s) - Q(s,a))^2]$$

#### 4.1 Create a QNetwork

Feed Forward network

```
[6]: fc_layer_params = (100,) # initial_collect_steps? input layer?

q_net = q_network.QNetwork(
    train_env.observation_spec(),
    train_env.action_spec(),
    # A list of fully_connected parameters,
    # where each item is the number of units in the layer
    fc_layer_params=fc_layer_params)
#q_net
```

#### 4.2 Instantiate a DqnAgent

Implements the DQN algorithm from

“Human level control through deep reinforcement learning” Mnih et al., 2015

<https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>

```
[7]: optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=learning_rate) # Adam algorithm

train_step_counter = tf.Variable(0)

agent = dqn_agent.DqnAgent(
    train_env.time_step_spec(),
    train_env.action_spec(),
```

```

q_network=q_net,
optimizer=optimizer, # AdamOptimizer function
td_errors_loss_fn=common.element_wise_squared_loss, # loss function
train_step_counter=train_step_counter) # integer step counter

agent.initialize()

```

### 4.3 Policies/Rules

The rules will return an action to produce the desired rewards

```

[8]: eval_policy = agent.policy # The main policy that is used for evaluation and
      ↪ deployment.
      collect_policy = agent.collect_policy # A second policy that is used for data
      ↪ collection.

```

## 5 Evaluation

Computes the average return of a policy per episode.

```

[9]: def compute_avg_return(environment, policy, num_episodes=10):

    total_return = 0.0
    for _ in range(num_episodes):

        time_step = environment.reset()
        episode_return = 0.0

        while not time_step.is_last():
            action_step = policy.action(time_step)
            time_step = environment.step(action_step.action)
            episode_return += time_step.reward
            total_return += episode_return

    avg_return = total_return / num_episodes
    return avg_return.numpy()[0]

```

Show baseline performance by random selection.

```

[10]: random_policy = random_tf_policy.RandomTFPolicy(train_env.time_step_spec(),
                                                       train_env.action_spec())
      compute_avg_return(eval_env, random_policy, num_eval_episodes)

```

```
[10]: 18.7
```

## 6 Data Collection

### 6.1 Replay Buffer

Dictionary?

```
[11]: replay_buffer = tf_uniform_replay_buffer.TFUniformReplayBuffer(  
    data_spec=agent.collect_data_spec,  
    batch_size=train_env.batch_size,  
    # The maximum number of items that can be stored in a  
    # single batch segment of the buffer  
    max_length=replay_buffer_max_length)  
  
agent.collect_data_spec, agent.collect_data_spec._fields  
  
[11]: (Trajectory(step_type=TensorSpec(shape=(), dtype=tf.int32, name='step_type'),  
    observation=BoundedTensorSpec(shape=(4,), dtype=tf.float32, name='observation',  
    minimum=array([-4.8000002e+00, -3.4028235e+38, -4.1887903e-01, -3.4028235e+38],  
        dtype=float32), maximum=array([4.8000002e+00, 3.4028235e+38,  
    4.1887903e-01, 3.4028235e+38],  
        dtype=float32)), action=BoundedTensorSpec(shape=(), dtype=tf.int64,  
    name='action', minimum=array(0), maximum=array(1)), policy_info=(),  
    next_step_type=TensorSpec(shape=(), dtype=tf.int32, name='step_type'),  
    reward=TensorSpec(shape=(), dtype=tf.float32, name='reward'),  
    discount=BoundedTensorSpec(shape=(), dtype=tf.float32, name='discount',  
    minimum=array(0., dtype=float32), maximum=array(1., dtype=float32))),  
    ('step_type',  
     'observation',  
     'action',  
     'policy_info',  
     'next_step_type',  
     'reward',  
     'discount'))
```

Recording the data in the replay buffer.

```
[12]: def collect_step(environment, policy, buffer):  
    time_step = environment.current_time_step()  
    action_step = policy.action(time_step)  
    next_time_step = environment.step(action_step.action)  
    traj = trajectory.from_transition(time_step, action_step, next_time_step)  
  
    # Add trajectory to the replay buffer  
    buffer.add_batch(traj)  
  
def collect_data(env, policy, buffer, steps):  
    for _ in range(steps):  
        collect_step(env, policy, buffer)
```

```
collect_data(train_env, random_policy, replay_buffer, initial_collect_steps)
```

Covert the dictionary as a data set and return an iterator.

```
[13]: dataset = replay_buffer.as_dataset(  
    num_parallel_calls=3,  
    sample_batch_size=batch_size,  
    num_steps=2).prefetch(3)  
    #dataset  
    iterator = iter(dataset)
```

WARNING:tensorflow:From /opt/anaconda3/lib/python3.8/site-packages/tensorflow/python/autograph/operators/control\_flow.py:1004: ReplayBuffer.get\_next (from tf\_agents.replay\_buffers.replay\_buffer) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use `as\_dataset(..., single\_deterministic\_pass=False)` instead.

## 7 Train

Two things must happen during the training loop:

collect data from the environment

use that data to train the agent's neural network(s)

```
[14]: try:  
    %%time  
except:  
    pass  
# (Optional) Optimize by wrapping some of the code in a graph using TF function.  
agent.train = common.function(agent.train)  
  
# Reset the train step  
agent.train_step_counter.assign(0)  
  
# Evaluate the agent's policy once before training.  
avg_return = compute_avg_return(eval_env, agent.policy, num_eval_episodes)  
returns = [avg_return]  
  
for _ in range(num_iterations):  
  
    # Collect a few steps using collect_policy and save to the replay buffer.  
    collect_data(train_env, agent.collect_policy, replay_buffer,   
    ↪collect_steps_per_iteration)  
  
    # Sample a batch of data from the buffer and update the agent's network.
```

```

experience, unused_info = next(iterator)
train_loss = agent.train(experience).loss

step = agent.train_step_counter.numpy()

if step % log_interval == 0:
    print('step = {0}: loss = {1}'.format(step, train_loss))

if step % eval_interval == 0:
    avg_return = compute_avg_return(eval_env, agent.policy, num_eval_episodes)
    print('step = {0}: Average Return = {1}'.format(step, avg_return))
    returns.append(avg_return)

```

WARNING:tensorflow:From /opt/anaconda3/lib/python3.8/site-packages/tensorflow/python/util/dispatch.py:201: calling foldr\_v2 (from tensorflow.python.ops.functional\_ops) with back\_prop=False is deprecated and will be removed in a future version.

Instructions for updating:

back\_prop=False is deprecated. Consider using tf.stop\_gradient instead.

Instead of:

```
results = tf.foldr(fn, elems, back_prop=False)
```

Use:

```
results = tf.nest.map_structure(tf.stop_gradient, tf.foldr(fn, elems))
```

```

step = 200: loss = 18.008861541748047
step = 400: loss = 5.086471080780029
step = 600: loss = 15.28239631652832
step = 800: loss = 9.468496322631836
step = 1000: loss = 11.550569534301758
step = 1000: Average Return = 15.699999809265137
step = 1200: loss = 35.28739929199219
step = 1400: loss = 2.0952863693237305
step = 1600: loss = 2.0557005405426025
step = 1800: loss = 25.886310577392578
step = 2000: loss = 5.963220119476318
step = 2000: Average Return = 24.299999237060547
step = 2200: loss = 11.770583152770996
step = 2400: loss = 8.902961730957031
step = 2600: loss = 37.822486877441406
step = 2800: loss = 6.025323867797852
step = 3000: loss = 9.689393043518066
step = 3000: Average Return = 110.4000015258789
step = 3200: loss = 22.36556625366211
step = 3400: loss = 32.306640625
step = 3600: loss = 15.72234058380127
step = 3800: loss = 6.93696403503418
step = 4000: loss = 50.098915100097656
step = 4000: Average Return = 90.0

```

step = 4200: loss = 47.52417755126953  
step = 4400: loss = 155.73275756835938  
step = 4600: loss = 153.7744140625  
step = 4800: loss = 112.13093566894531  
step = 5000: loss = 15.244081497192383  
step = 5000: Average Return = 114.5999984741211  
step = 5200: loss = 21.695892333984375  
step = 5400: loss = 5.562005996704102  
step = 5600: loss = 6.435539245605469  
step = 5800: loss = 59.24911117553711  
step = 6000: loss = 7.506043910980225  
step = 6000: Average Return = 160.0  
step = 6200: loss = 65.04222869873047  
step = 6400: loss = 217.68240356445312  
step = 6600: loss = 174.95286560058594  
step = 6800: loss = 164.0635528564453  
step = 7000: loss = 10.535503387451172  
step = 7000: Average Return = 114.69999694824219  
step = 7200: loss = 117.93932342529297  
step = 7400: loss = 67.07676696777344  
step = 7600: loss = 10.39604377746582  
step = 7800: loss = 67.12391662597656  
step = 8000: loss = 105.36053466796875  
step = 8000: Average Return = 198.5  
step = 8200: loss = 118.29344177246094  
step = 8400: loss = 140.25460815429688  
step = 8600: loss = 30.344758987426758  
step = 8800: loss = 11.699833869934082  
step = 9000: loss = 476.17034912109375  
step = 9000: Average Return = 176.1999969482422  
step = 9200: loss = 451.31793212890625  
step = 9400: loss = 20.707763671875  
step = 9600: loss = 28.288009643554688  
step = 9800: loss = 125.00628662109375  
step = 10000: loss = 165.96885681152344  
step = 10000: Average Return = 164.39999389648438  
step = 10200: loss = 234.79196166992188  
step = 10400: loss = 34.319217681884766  
step = 10600: loss = 73.28331756591797  
step = 10800: loss = 21.278406143188477  
step = 11000: loss = 207.4669189453125  
step = 11000: Average Return = 196.5  
step = 11200: loss = 9.321070671081543  
step = 11400: loss = 34.81235885620117  
step = 11600: loss = 24.20108985900879  
step = 11800: loss = 218.18382263183594  
step = 12000: loss = 54.903507232666016  
step = 12000: Average Return = 200.0

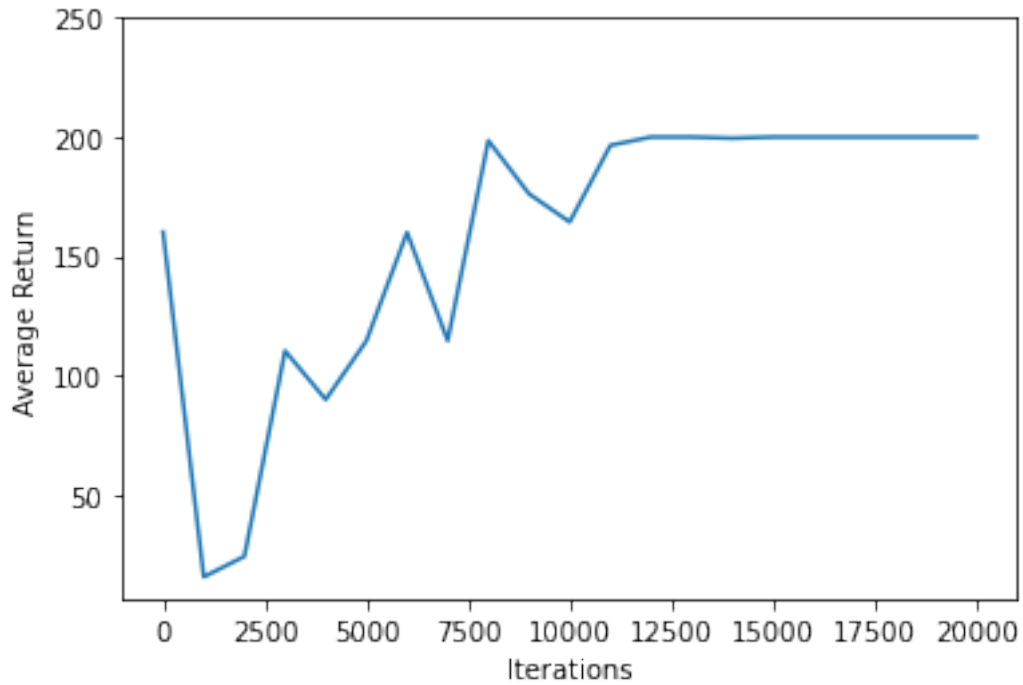
step = 12200: loss = 696.0888061523438  
step = 12400: loss = 118.06753540039062  
step = 12600: loss = 29.727664947509766  
step = 12800: loss = 172.8981475830078  
step = 13000: loss = 36.71021270751953  
step = 13000: Average Return = 200.0  
step = 13200: loss = 43.96297836303711  
step = 13400: loss = 187.39060974121094  
step = 13600: loss = 699.5609741210938  
step = 13800: loss = 43.06932830810547  
step = 14000: loss = 262.7250061035156  
step = 14000: Average Return = 199.5  
step = 14200: loss = 44.86981201171875  
step = 14400: loss = 788.365966796875  
step = 14600: loss = 186.86534118652344  
step = 14800: loss = 34.16826248168945  
step = 15000: loss = 68.5713882446289  
step = 15000: Average Return = 200.0  
step = 15200: loss = 36.6361083984375  
step = 15400: loss = 67.6352767944336  
step = 15600: loss = 88.42098999023438  
step = 15800: loss = 66.49943542480469  
step = 16000: loss = 34.3782958984375  
step = 16000: Average Return = 200.0  
step = 16200: loss = 51.016273498535156  
step = 16400: loss = 979.0067138671875  
step = 16600: loss = 20.1558780670166  
step = 16800: loss = 74.05459594726562  
step = 17000: loss = 27.95159912109375  
step = 17000: Average Return = 200.0  
step = 17200: loss = 149.70065307617188  
step = 17400: loss = 16.747379302978516  
step = 17600: loss = 678.1433715820312  
step = 17800: loss = 41.89617919921875  
step = 18000: loss = 609.5775756835938  
step = 18000: Average Return = 200.0  
step = 18200: loss = 454.809326171875  
step = 18400: loss = 48.18207550048828  
step = 18600: loss = 48.29536437988281  
step = 18800: loss = 122.12271118164062  
step = 19000: loss = 1473.701416015625  
step = 19000: Average Return = 200.0  
step = 19200: loss = 135.0766143798828  
step = 19400: loss = 408.56280517578125  
step = 19600: loss = 489.67938232421875  
step = 19800: loss = 107.77105712890625  
step = 20000: loss = 753.5126953125  
step = 20000: Average Return = 200.0



## 8 Visualization

```
[15]: iterations = range(0, num_iterations + 1, eval_interval)
plt.plot(iterations, returns)
plt.ylabel('Average Return')
plt.xlabel('Iterations')
plt.ylim(top=250)
```

```
[15]: (6.484999799728394, 250.0)
```



## 9 Reference

[https://www.tensorflow.org/agents/tutorials/1\\_dqn\\_tutorial](https://www.tensorflow.org/agents/tutorials/1_dqn_tutorial)